
Subgoal Search For Complex Reasoning Tasks

Konrad Czechowski *
University of Warsaw

Tomasz Odrzygóźdź*
University of Warsaw

Marek Zbysiński
University of Warsaw

Michał Zawalski
University of Warsaw

Krzysztof Olejnik
University of Warsaw

Yuhuai Wu
University of Toronto

Łukasz Kuciński
Polish Academy of Sciences

Piotr Miłoś
Polish Academy of Sciences

Abstract

Humans excel in solving complex reasoning tasks through a mental process of moving from one idea to a related one. Inspired by this, we propose Subgoal Search (kSubS) method. Its key component is a learned subgoal generator, an analog of human intuition, that produces a diversity of subgoals that are both achievable and closer to the solution. Using subgoals reduces the search space and induces a high-level search graph suitable for efficient planning. In this paper, we implement kSubS using a transformer-based subgoal module coupled with the classical best-first search framework. We show that a simple approach of generating k -the step ahead subgoals is surprisingly efficient on three challenging domains: two popular puzzle games, Sokoban and the Rubik’s Cube, and an inequality proving benchmark INT. kSubS achieves strong results including state-of-the-art on INT within a modest computational budget.

1 Introduction

Reasoning is often regarded as a defining property of advanced intelligence [30, 14]. When confronted with a complicated task, humans’ thinking process often moves from one idea to a related idea, and the progress is made through milestones, or *subgoals*, rather than through atomic actions that are necessary to transition between subgoals [12]. During this process, thinking about one subgoal can lead to a possibly diverse set of subsequent subgoals that are conceptually reachable and make a promising step towards the problem’s solution. This intuitive introspection is backed by neuroscience evidence [15], and in this work, we present an algorithm that mimics this process. Our approach couples a deep learning generative subgoal modeling with classical search algorithms to allow for successful planning with subgoals. We showcase the efficiency of our method on the following complex reasoning tasks: two popular puzzle games Sokoban and the Rubik’s Cube, and an inequality theorem proving benchmark INT [46], achieving the state-of-the-art results in INT and competitive results for the remaining two.

The deep learning revolution has brought spectacular advancements in pattern recognition techniques and models. Given the hard nature of reasoning problems, these are natural candidates to provide search heuristics [4]. Indeed, such a blend can produce impressive results [35, 36, 29, 1]. These

*equal contribution

approaches seek solutions using elementary actions. Others, e.g. [23, 26, 18], utilize variational subgoals generators to deal with long-horizon visual tasks. We show that these ideas can be pushed further to provide algorithms capable of dealing with combinatorial complexity.

We present Subgoal Search (SubS) method and give its practical implementations: MCTS-kSubS and BF-kSubS. SubS consists of the following four components: planner, subgoal generator, a low-level policy, and a value function. The planner is used to search over the graph induced by the subgoal generator and is guided by the value function. The role of the low-level policy is to prune the search tree as well as to transition between subgoals. In this paper, we assume that the generator predicts subgoals that are k step ahead (towards the solution) from the current state, and to emphasize this we henceforth add k to the method’s abbreviation. MCTS-kSubS and BF-kSubS differ in the choice of the search engine: the former uses Monte-Carlo Tree Search (MCTS), while the latter is backed by Best-First Search (BestFS). We provide two sets of implementations for the generator, the low-level policy, and the value functions. The first one uses transformer architecture ([40]) for each component, while the second utilizes a convolutional network for the generator and the value function, and the classical breadth-first search for the low-level policy. This lets us showcase the versatility and effectiveness of the approach.

The subgoal generator lies at the very heart of Subgoal Search, being an analog of human intuition and implementation of reasoning with high-level ideas. To be useful in a broad spectrum of contexts, the generator should be implemented as a learnable (generative) model. As a result, it is expected to be imperfect and (sometimes) generate incorrect predictions, which may turn the search procedure invalid. Can we thus make planning with learned subgoals work? In this paper, we answer this question affirmatively: we show that the autoregressive framework of transformer-based neural network architecture [41] leads to superior results in challenging domains.

We train the transformer with the objective to predict the k -th step ahead. The main advantages of this subgoal objective are simplicity and empirical efficiency. We used expert data to generate labels for supervised training. When offline datasets are available, which is the case for the environments considered in this paper², such an approach allows for stable and efficient optimization with high-quality gradients. Consequently, this method is often taken when dealing with complex domains (see e.g. [33, 43]) or when only an offline expert is available³. Furthermore, we found evidence of out-of-distribution generalization.

Finally, we formulate the following hypothesis aiming to shed some light on why kSubS is successful: we speculate that subgoal generation may alleviate errors in the value function estimation. Planning methods based on learning, including kSubS, typically use imperfect value function-based information to guide the search. While traditional low-level search methods are susceptible to local noise, subgoal generation allows for evaluations of the value functions at temporally distant subgoals, which improves the signal-to-noise ratio and allows a “leap over” the noise.

To sum up, our contributions are:

1. We propose Subgoal Search method with two transformer-based implementations (MCTS-kSubS, BF-kSubS), and one, simpler, convolutional-based implementation. We demonstrate that our approach requires a relatively little search or, equivalently, is able to handle bigger problems. Interestingly, we observe evidence of out-of-distribution generalization.
2. We show that a transformer-based autoregressive model learned with a simple supervised objective to produce k -th step ahead subgoal already achieves a strong performance.
3. We study the negative influence of value function errors on planning and show that using subgoal planning might mitigate this problem.

We provide the code of our method and experiment settings at <https://github.com/subgoal-search/subgoal-search>.

²The dataset for INT or Sokoban can be easily generated or are publicly available. For the Rubik’s Cube, we use random data or simple heuristic (random data are often sufficient for robotic tasks and navigation.)

³For example, the INT engine can easily generate multiple proves of random statements, but *cannot* prove a given theorem.

2 Related work

In classical AI, reasoning is often achieved by *search* ([30]). Search rarely can be exhaustive and a large body of algorithms and heuristics has been developed over the years, [30, Section 3.5]. It is hypothesized that progress can be achieved by combining search with learning [4]. Among notable successful examples of such a blend are Alpha Zero [34], or solving Rubik’s cube using a learned function to guide A* search (see [1]).

An eminent early example of using goal-directed reasoning is PARADISE algorithm ([44]). In deep learning literature, [19] was perhaps the first work implementing subgoal planning. This was followed by a large body of work on planning with subgoals in the latent space for visual tasks ([18, 24, 26, 23, 16, 8, 27]) or landmark-based navigation methods ([31], [20], [11], [37], [47]).

The tasks considered in the aforementioned studies are often quite forgiving when it comes to small errors in the subgoal generation. This can be contrasted with complex reasoning domains, in which small changes to a given state may drastically change its meaning or even render it invalid. Thus, neural networks may struggle to generate semantically valid states ([4, Section 6.1]). Assuming that this problem was solved, a generated subgoal still remains to be assessed. The exact evaluation may, in general, require exhaustive search or access to an oracle (in which case the original problem is essentially solved). Consequently, it is unlikely that a simple planner (e.g., one unrolling independent sequences of subgoals [8]) will either produce an informative outcome or could be easily improved using only local changes via gradient descent [24], or cross-entropy method (CEM) [23, 26, 27]. Existing approaches based on more powerful subgoal search methods, on the other hand, have their limitations. [10] is perhaps the closest to our method and uses an MCTS to search the subgoal-induced graph. However, it uses a predefined (not learned) predicate function as a subgoal generator, limiting applicability to the problems with available high-quality heuristics. Learned subgoals are used in [25], a method of hierarchical planning of subgoal sequences. That said, the subgoal space needs to be relatively small for this method to work (or crafted manually to reduce cardinality).

More generally, concepts related to goals and subgoals percolated to reinforcement learning early on, leading, among others, to prominent ideas like hindsight [17], hierarchical learning [39, 6] or the Horde architecture [38]. Recently, with the advent of *deep* reinforcement learning, these ideas have been resurfacing and scaled up to deliver their initial promises. For example, [42] implements ideas of [6] and a very successful hindsight technique [3] is already considered to be a core RL method. Further, a recent paper [28] utilizes a maximum entropy objective to select achievable goals in hindsight training.

3 Method

Our method is designed for problems, which can be formulated as a search over a graph with a known transition model. Formally, let G be a directed graph and $\tilde{S} \subset G$ be the set of success states. We assume that, during the solving process, the algorithm can, for a given node g , determine the edges starting at g and check if $g \in \tilde{S}$.

Subgoal Search method consists of four components: planner, subgoal generator, low-level policy, and a value function. The planner, coupled with a value function, is used to search over the graph induced by the subgoal generator. Namely, for each selected subgoal, the generator allows for sampling the candidates for the next subgoals. Only these reachable by the low-level policy are used. The procedure continues until the solution is found or the computational budget is exhausted.

In this paper we provide BestFS- and MCTS- backed implementations of kSubS. Algorithm 1 presents BF-kSubS (see Algorithm 4 in Appendix A.1 for MCTS-kSubS), while the listing for low-level policy and the subgoal generator is given in Algorithms 2 and Algorithm 3, respectively.

For INT and Rubik’s Cube, we use transformer models (see Appendix B.1) in all components other than the planner. The setting for Sokoban is different (based on convnets), and for clarity, we postpone its description till the end of this section. While transformers could also be used in Sokoban, we show that a simplified setup already achieves strong results. This showcases that Subgoal Search is general enough to work with different design choices. For each environment we collect expert data, see Section 4.2, on which learned models are trained.

Subgoal generator. Formally, it is a mapping $\rho: \mathcal{S} \rightarrow \mathbf{P}(\mathcal{S})$, where $\mathbf{P}(\mathcal{S})$ is a family of probability distributions over the environment’s state space \mathcal{S} . More precisely, let us define a trajectory as a sequence of state-action pairs $(s_0, a_0), (s_1, a_1), \dots, (s_n, a_n)$, with $(s_i, a_i) \in \mathcal{S} \times \mathcal{A}$, where \mathcal{A} stands for the action space and n is the trajectory’s length. We will say that the generator predicts k -step ahead subgoals, if at any state s_ℓ it aims to predict $s_{\min(\ell+k, n)}$. We show, perhaps surprisingly, that this simple objective is an efficient way to improve planning, even for small values of k , i.e. $k \in \{2, 3, 4, 5\}$.

Operationally, the subgoal generator takes as input an element of the \mathcal{S} and returns *subgoals*, a set of new candidate states expected to be closer to the solution, and is implemented by Algorithm 3. This procedure is parameterized by three constants: C_3 , governing the number of subgoals to be generated, C_4 , the number of beams in beam search, and C_5 , a probability threshold which controls how plausible (according to the model) the subgoal set is. In INT and Rubik’s Cube, we model subgoal generator network (ρ in Algorithm 3) as transformer. It is trained in a supervised way on the expert data, with training examples being: s_ℓ (input) and $s_{\min(\ell+k, n)}$ (label), see Appendix D for details. Other possible approaches to modelling a subgoal generator are discussed in Section 5.

Algorithm 1 Best-First Subgoal Search

Require: C_1 max number of nodes
 V value function network
 SOLVED predicate of solution

function SOLVE(s_0)
 $T.\text{PUSH}((V(s_0), s_0))$ $\triangleright T$ is priority queue
 $\text{paths}[s_0] \leftarrow []$ $\triangleright \text{paths}$ is dict of lists
 $\text{seen}.\text{ADD}(s_0)$ $\triangleright \text{seen}$ is set
 while $0 < \text{LEN}(T)$ and $\text{LEN}(\text{seen}) < C_1$ **do**
 $s \leftarrow T.\text{EXTRACT_MAX}()$
 $\text{subgoals} \leftarrow \text{SUB_GENERATE}(s)$
 for s' **in** subgoals **do**
 if s' **in** seen **then** **continue**
 $\text{seen}.\text{ADD}(s')$
 $\text{actions} \leftarrow \text{GET_PATH}(s, s')$
 if $\text{actions}.\text{EMPTY}()$ **then**
 continue
 $T.\text{PUSH}((V(s'), s'))$
 $\text{paths}[s'] \leftarrow \text{paths}[s] + \text{actions}$
 if $\text{SOLVED}(s')$ **then**
 return $\text{paths}[s']$

return False

Algorithm 2 Low-level conditional policy

Require: C_2 steps limit
 π low-level conditional policy network
 M model of the environment

function GET_PATH(s_0 , subgoal)
 $\text{step} \leftarrow 0$
 $s \leftarrow s_0$
 $\text{action_path} \leftarrow []$
 while $\text{step} < C_2$ **do**
 $\text{action} \leftarrow \pi.\text{PREDICT}(s, \text{subgoal})$
 $\text{action_path}.\text{APPEND}(\text{action})$
 $s \leftarrow M.\text{NEXT_STATE}(s, \text{action})$
 if $s = \text{subgoal}$ **then**
 return action_path
 $\text{step} \leftarrow \text{step} + 1$
 return $[]$

Low-level (conditional) policy. Formally, it is a mapping $\pi: \mathcal{S} \times \mathcal{S} \rightarrow \mathcal{A}^*$. It provides a sequence of actions on how to reach a subgoal starting from a given initial state. Operationally, it may return an empty sequence if the subgoal cannot be reached within C_2 steps, see Algorithm 2. This is used as a pruning mechanism for the *subgoals* set in Algorithm 1.

In INT and Rubik’s Cube, we implement low-level policy network (π in Algorithm 2) as a transformer (see Appendix B.1 for details). Similarly to the subgoal generator, it is trained using expert data in a supervised fashion, i.e. for a pair $(s_\ell, s_{\min(\ell+i, n)})$, with $i \leq k$, its objective is to predict a_ℓ .

Value function. Formally, it is a mapping $V: \mathcal{S} \rightarrow \mathbb{R}$, that assigns to each state a value related to its distance to the solution, and it is used to guide the search (see Algorithm 1 and Algorithm ??). For INT and Rubik’s Cube, we use a transformer architecture, which is trained on expert data, to predict for each state s_ℓ its negated distance to the goal $\ell - n$, where n denotes the end step of a trajectory that s_ℓ belongs to.

Planner. This is the engine that we use to search the subgoal-induced graph. In this paper, we use BestFS (Algorithm 1) and MCTS (Algorithm ?? in Appendix A.1). The former is a classic planning method, which maintains a priority queue of states waiting to be explored, and greedily (with respect to their value) selects elements from it (see, e.g., [30]). MCTS is a search method that iteratively and

explicitly builds a search tree, using (and updating) the collected node statistics (see, e.g., [5]). In this paper, we use an AlphaZero-like [33] algorithm for single-player games.

Sokoban setup. Thanks to a low branching factor and small values of k that we consider in this paper, the setup for Sokoban can be simplified. For the neural network used by the subgoal generator (ρ in Algorithm 3) we use a convolutional-network based architecture. Accordingly, we replace the beam search of Algorithm 4 with an alternative approach, better suited for the chosen architecture, see Appendix B.2. The value function used in this implementation comes from a separate MCTS training, for details see Appendix C.1. For a low-level policy, we simply use breadth-first search, instead of Algorithm 2.

Algorithm 3 Subgoal generator

Require: C_3 number of subgoals
 C_4 number of beams in sampling
 C_5 target probability
 ρ subgoal generator network

function SUB_GENERATE(s)
 $\text{subgoals} \leftarrow \emptyset$
 $\text{states}, \text{probs} \leftarrow \text{BEAM_SEARCH}(\rho, s; C_3, C_4)$
 $\triangleright (\text{states}, \text{probs})$ is sorted wrt probs
 $\text{total_p} \leftarrow 0$
 for $\text{state}, p \in (\text{states}, \text{probs})$ **do**
 if $\text{total_p} > C_5$ **then break**
 $\text{subgoals.ADD}(\text{state})$
 $\text{total_p} \leftarrow \text{total_p} + p$
 return subgoals

4 Experiments

In this section, we empirically demonstrate the efficiency of MCTS-kSubS and BF-kSubS, in particular, we show that they vastly outperform their standard (“non-subgoal”) counterparts. As a test ground, we consider three challenging domains: Sokoban, Rubik’s Cube, and INT. All of them, require non-trivial reasoning. The Rubik’s Cube is a well-known 3D combination puzzle. Sokoban is a complex video puzzle game known to be NP-hard and thus a challenging testbed for planning methods. INT [46] is a recent theorem proving benchmark.

4.1 Training protocol and baselines

Our experiments consist of three stages. First, we collect domain-specific expert data, see Section 4.2. Secondly, we train the subgoal generator, low-level conditional policy, and value function networks using the data and targets described in Section 3. For more details see Appendix D. Eventually, we evaluate the planning performance of MCTS-kSubS and BF-kSubS, details of which are presented below. In the second step, we use supervised learning, which makes our setup stable with respect to network initialization, see details in Appendix D.1.3.

As baselines, we use BestFS and MCTS (single player reimplement of AlphaZero) operating on elementary actions. In INT and Rubik’s Cube both algorithms utilize policy network (trained with behavioral cloning, on the same dataset, which we used to train kSubS). Note that policy network can be seen as a variant of subgoal generator for $k = 1$, since we consider deterministic transition function between states (the distribution over actions can be viewed as a distribution over states). This representation allowed us to reuse the same codebase for baseline search methods and kSubS and ensures a fair comparison. For more details see Appendix H.

4.2 Search Domains and Datasets

Sokoban is a single-player complex game in which a player controls an agent that is able to push boxes across the board. The goal is to place all boxes on target locations solely by pushing them, without crossing any obstacles or walls. Sokoban has recently been used to test the boundaries in RL [13]. Sokoban is known to be hard [9], mainly due to its combinatorial complexity and the existence of irreversible states. Formally, deciding if a given Sokoban board is solvable is a PSPACE-complete problem. [7]

To collect the expert dataset, we use a separate training using an MCTS agent [22] (with implementation made publicly available by the authors). Our dataset consists of all successful trajectories occurring during the training. These are suboptimal, especially in the early phases of the training or for harder boards.

Rubik’s Cube is a classical 3-dimensional classic puzzle. It is considered challenging due to the fact that the search space has more than 4.3×10^{18} possible configurations. Similarly to Sokoban, Rubik’s Cube has been recently used as a testing ground for RL methods [1].

To collect the expert dataset, we generate random paths of length 30 starting from the solved cube and take them backward. These backward solutions are highly sub-optimal (optimal solutions are proven to be shorter than 26).

INT: Inequality Benchmark for Theorem Proving. INT provides a generator of inequalities, which produces mathematical formulas along with a corresponding proof. Proofs are represented as sequences of consecutively applied mathematical axioms (there is a total number of 18 axioms). Each action is represented as a tuple containing a chosen axiom and its input entities. Action space in this problem can be very large, reaching up to 10^6 elements, which significantly complicates planning.

INT is equipped with the generator of mathematical proofs [46, Section 3.3]. It generates a path by randomly applying axioms starting with a trivial statement. The path taken backward constitutes the proof of its final statement. The path length is an important hyperparameter regulating the difficulty – we use 5, 10, 15. The proofs generated as follows are not guaranteed to be optimal.

Hyperparameters of BF-kSubS used in specific domains are shown in Table 1. For more details on datasets see Appendix C.

	INT	Sokoban	Rubik
k	3	4	4
C_1	400	5000	1500
C_2	4	4	7
C_3	4	4	3
C_4	16	NA	32
C_5	1	0.98	1

Table 1: BF-kSubS hyperparameters.

4.3 Main results

In this section, we present our most important finding: that MCTS-kSubS and BF-kSubS perform significantly better than the respective methods not using subgoals, including state-of-the-art on INT. Using Subgoal Search enables for more efficient search and consequently scales up to problems with higher difficulty.

In Figure 1, we present the performance of Subgoal Search. We measure the *success rate* in function of the *search budget* used. The success rate metric is measured on 1000 instances of the given problem (which results in confidence intervals within ± 0.03 range). For BF-kSubS the search budget is referred to as *graph size* and includes the total number of nodes visited by the algorithm. More precisely, these are subgoals visited in Algorithm 1 and nodes on paths connecting them (given by Algorithm 2). For MCTS, we report *MCTS passes*, see details in Appendix A.1.

Below we discuss the results separately for each domain.

INT. The difficulty of the problems in INT increases fast with the proof length (denoted by L in [46]) and the number of accessible axioms (denoted by K in [46]). We used $K = 18$; all of available axioms. We observed, that BF-kSubS scales to proofs of length 10 and 15, which are significantly harder than 5 considered in [46], see Table 2. The same holds for MCTS-kSubS, see Appendix A.2. We check also that MCTS-kSubS vastly outperforms the AlphaZero algorithm, see Figure 1 (top, left).⁴

An MCTS agent was also evaluated on INT by [46]. The Implementation provided there, guided by the value function realized by Graph Neural Network, obtained 92% solve rate on the proofs of length 5. We ensured that our baseline is stronger - it solves over 99% of the problems on the same problem set.

Sokoban. Using BF-kSubS allows for significantly higher success rates within the same computational budget, see Table ?? . Our solution works well for 12×12 (with 4 boxes) boards commonly used in recent deep-learning research [13, 2] but also scales to much harder boards - 16×16 and 20×20 . Importantly, we observe that already for small computational budget (graph size 100). BF-kSubS obtains higher success rates than the expert we used to create the dataset (these are 78%, 67%, 60% for the respective board sizes).

⁴We use our implementation of AlphaZero. In particular, it obtains stronger results than ones reported in [46] for $L = 5$. ($L = 10$ and larger are not considered in [46].)

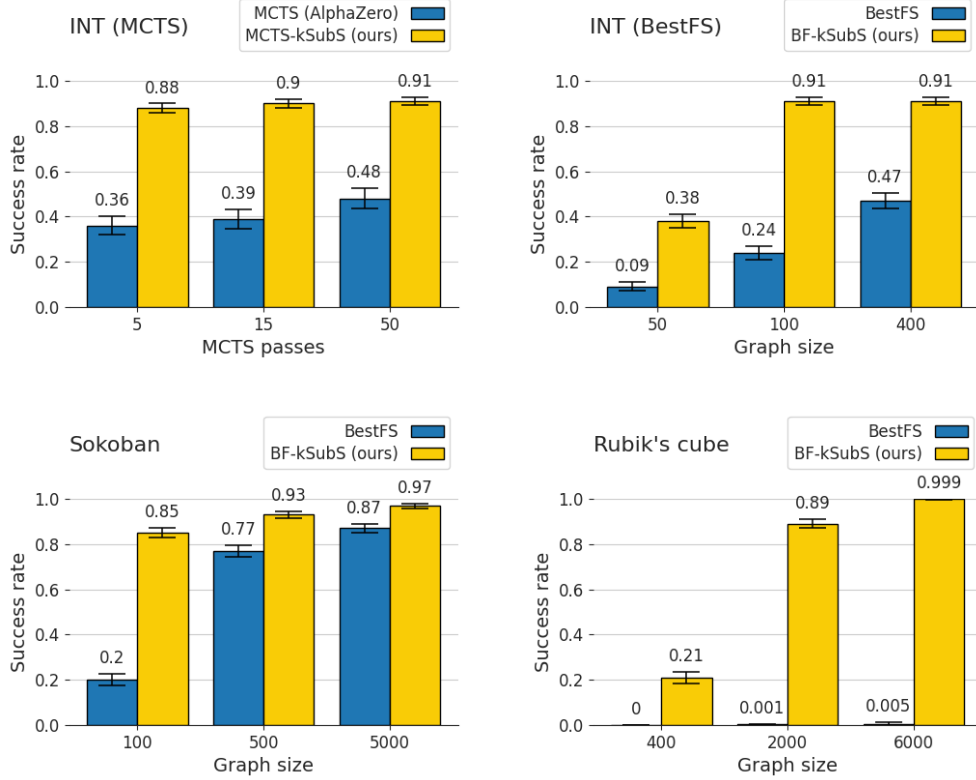


Figure 1: The performance of Subgoal Search. (top, left) compared on INT (with the proof length 15) to AlphaZero. (top, right) BF-kSubS consistently achieves high performance, even for small computational budgets. (bottom, left) similarly on Sokoban (board size 12x12 with 4 boxes) the advantage of BF-kSubS is prominent for small budget. (bottom, right) BestFS fails to solve Rubik’s Cube.

Proof length		5		10		15	
Method		BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)
Graph size	50	0.82	0.99	0.47	0.97	0.09	0.38
	100	0.89	0.99	0.64	0.99	0.24	0.91
	200	0.92	0.99	0.67	0.99	0.35	0.91
	400	0.93	0.99	0.72	0.99	0.47	0.91

Table 2: INT success rates for various problems difficulties and graphs sizes.

Rubik’s Cube BF-kSubS solves nearly 100% of cubes, BestFS is close to 0%, see Figure 1 (bottom, right). This is perhaps the most striking example of the advantage of using subgoal generator instead of low-level actions. We present possible explanation in Appendix J. We also include there the comparison with [1], a recent work on the Rubik’s Cube, we note however their scope is to obtain nearly optimal solutions using high computational budgets.

Out-of-distribution generalization is considered to be the crucial ability to make progress in hard combinatorial optimization problems [4] and automated theorem proving [46]. The INT benchmark has been specifically designed to benchmark this phenomenon. We checked how the Subgoal Search, method trained on proofs on length 10, generalizes favourably to longer problems, see Figure 3.

It might be sometimes hard to compare computational budgets between various algorithms and their versions. In Appendix E we measure that BF-kSubS and MCTS-kSubS offer very practical benefits in terms of wall-time, sometimes as much as 7 times.

We provide examples of solutions and generated subgoals in Appendix G. The preceding sections are devoted to a more detailed analysis of various aspects of our method.

Board Size		12 x 12		16 x 16		20 x 20	
Method		BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)
Graph size	50	0.02	0.69	0.01	0.47	0.00	0.31
	100	0.20	0.85	0.10	0.74	0.04	0.64
	1000	0.81	0.95	0.80	0.89	0.73	0.88
	5000	0.87	0.97	0.89	0.93	0.87	0.91

Table 3: Sokoban success rates for various board sizes (each with 4 boxes).

4.4 Analysis of k (subgoal distance) parameter

The subgoals are trained to be k steps ahead. In principle, having higher k should make planning easier. However, as k increases, the quality of the generator might drop; thus, the overall effect is uncertain. For Sokoban increasing k generally helps, but the benefits saturate around $k = 4$, see Figure 2. We note that higher k increases computational demand from low-level conditional policy search.




Figure 2: Comparison of success rates w.r.t k .




Figure 3: Out-of-distribution generalization to longer proofs. We compare with the Behavioral Cloning agent (Policy) studied in [46].

4.5 Quality of subgoals

The learned subgoal generator is likely to be imperfect (especially in hard problems). We study this on smaller 10×10 boards of Sokoban, for which the true distance $dist$ to the solution can be found with the Dijkstra algorithm. In Figure 4, we study $\Delta := dist_{s_1} - dist_{s_2}$, where s_1 is a sampled state and s_2 is a subgoal generated from s_1 . Ideally, the histogram should concentrate on $k = 4$ used in training. We see, however, that in more than 60% subgoals lead to an improvement, see Figure 4.

We checked that in INT and Rubik’s Cube about 50% of generated subgoals can be reached by the low-level conditional policy in Algorithm 2.




Figure 4: Histogram of Δ . Note that 12% of subgoals increases the distance, with 5% leading to irreversible “dead states” present in Sokoban.

4.6 Value errors

There are a few possible explanations for the success of our method. Below we discuss one of them. We hypothesize that using subgoals mitigates errors of value estimates, inherently present when using learned models. Both BestFS and MCTS choose where to expand search graph (or tree) based on the value function predictions. If on the path leading to the solution, value estimates are not monotonic, it will likely slow down the process of approaching the goal state. Intuitively, if we consider a path in higher temporal resolution i.e. one with consecutive states spaced k actions apart, the value estimates

become more monotonic. We studied this on proofs generated by INT, the probability that value decreases when moving ℓ steps changes from 0.32 when $\ell = 1$ down to only 0.02 for $\ell = 4$.

Another problem is over-optimism, a search method would be misguided into states with erroneously positive values. To illustrate this consider $S(s)$, the set of all states within distance 3 from s and, additionally, having the same distance to the solution as s . Intuitively, $S(s)$ are similar states with respect to the difficulty of solving. Unlikely, the variance of value function prediction for $S(s)$ is equal to 0.84 (averaged over different s). For comparison, the average increase of value for moving one step closer to the solution is only 0.33. We found that with probability 86% there exists a suboptimal state in $S(s)$, i.e. having a higher value than the best immediate neighbour of s (which by properties of the game will be closer to solution in Sokoban). If instead of the neighbours we consider states closer by 4 step (say given by a subgoal generator), then the probability of the error mentioned above drops to 38%. For more details see Appendix F.2.

5 Limitations and future work

In this section, we list some limitations of our work and suggest further research directions.

Reliance on expert data In this version, we use expert data to train learnable models. As kSubS improves the performance, we speculate that training akin to AlphaZero can be used, i.e. in a planner-learner loop without any outside knowledge.

Optimality and completeness kSubS searches over a reduced state space, which might produce suboptimal solutions or even fail to find them. This is arguably unavoidable if we seek an efficient method for complex problems.

Subgoals definition We use simple k -step ahead subgoals, which is perhaps not always optimal. Our method can be coupled with other subgoal paradigms. Unsupervised detection of landmarks (see e.g. [47]) seems an attractive option.

More environment In future work, we plan to test kSubS on more environments to understand its strengths and weaknesses better. In this work, we generate subgoals in the state space, which might be limiting for tasks with high dimensional input (e.g., visual).

Reliance on a model of the environment We use a perfect model of the environment, which is a common practice for some environments, e.g., INT. Extending kSubS to use learned (imperfect) models is an important future research direction.

Determinism Our method requires the environment to be deterministic.

6 Conclusions

We propose Subgoal Search, a search algorithm based on subgoal generator. We present two practical implementations MCTS-kSubS and BF-kSubS meant to be effective in complex domains requiring reasoning. We confirm that indeed our implementations excel in Sokoban, Rubik’s Cube, and inequality benchmark INT. Interestingly, a simple k step ahead mechanism of generating subgoals backed up by transformer-based architectures perform surprisingly well. This evidence, let us hypothesise, that our methods (and related) can be further scaled up to even harder reasoning tasks.

6.1 Potential Negative Societal Impacts

We investigate a general-purpose search algorithm. The progress in domains we use as testing grounds (Automated Theorem Proving, combinatorially complex puzzles) arguably has relatively low societal risk. One of the potential applications of ATP in real-world is to prove the correctness of software [32], which could be used both to fix or to compromise it.

Our method can be potentially used to improve algorithms in various domains, including SAT-solving, chemical retrosynthesis planning, Combinatorial Optimization, and automated code synthesis. Analysis of the potential impacts of development in these fields is out of the scope of this paper.

References

- [1] Forest Agostinelli, Stephen McAleer, Alexander Shmakov, and Pierre Baldi. Solving the rubik’s cube with deep reinforcement learning and search. *Nature Machine Intelligence*, 1(8):356–363, 2019.
- [2] Ilge Akkaya, Marcin Andrychowicz, Maciek Chociej, Mateusz Litwin, Bob McGrew, Arthur Petron, Alex Paino, Matthias Plappert, Glenn Powell, Raphael Ribas, et al. Solving rubik’s cube with a robot hand. *arXiv preprint arXiv:1910.07113*, 2019.
- [3] Marcin Andrychowicz, Dwight Crow, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, Pieter Abbeel, and Wojciech Zaremba. Hindsight experience replay. In *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, 4-9 December 2017, Long Beach, CA, USA*, pages 5048–5058, 2017.
- [4] Yoshua Bengio, Andrea Lodi, and Antoine Prouvost. Machine learning for combinatorial optimization: a methodological tour d’horizon. *European Journal of Operational Research*, 2020.
- [5] Cameron B Browne, Edward Powley, Daniel Whitehouse, Simon M Lucas, Peter I Cowling, Philipp Rohlfshagen, Stephen Tavenor, Diego Perez, Spyridon Samothrakis, and Simon Colton. A survey of monte carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in games*, 4(1):1–43, 2012.
- [6] Peter Dayan and Geoffrey E. Hinton. Feudal reinforcement learning. In Stephen Jose Hanson, Jack D. Cowan, and C. Lee Giles, editors, *Advances in Neural Information Processing Systems 5, [NIPS Conference, Denver, Colorado, USA, November 30 - December 3, 1992]*, pages 271–278. Morgan Kaufmann, 1992.
- [7] Dorit Dor and Uri Zwick. Sokoban and other motion planning problems. *Computational Geometry*, 13(4):215–228, 1999.
- [8] Kuan Fang, Yuke Zhu, Animesh Garg, Silvio Savarese, and Li Fei-Fei. Dynamics learning with cascaded variational inference for multi-step manipulation. *arXiv preprint arXiv:1910.13395*, 2019.
- [9] Alan Fern, Roni Khordon, and Prasad Tadepalli. The first learning track of the international planning competition. *Machine Learning*, 84(1-2):81–107, 2011.
- [10] Thomas Gabor, Jan Peter, Thomy Phan, Christian Meyer, and Claudia Linnhoff-Popien. Subgoal-based temporal abstraction in Monte-Carlo Tree Search. In *IJCAI*, pages 5562–5568, 2019.
- [11] Wei Gao, David Hsu, Wee Sun Lee, Shengmei Shen, and Karthikk Subramanian. Intention-net: Integrating planning and deep learning for goal-directed autonomous navigation. In *1st Annual Conference on Robot Learning, CoRL 2017, Mountain View, California, USA, November 13-15, 2017, Proceedings*, volume 78 of *Proceedings of Machine Learning Research*, pages 185–194. PMLR, 2017.
- [12] Timothy Gowers. *The importance of mathematics*. Springer-Verlag, 2000.
- [13] Arthur Guez, Mehdi Mirza, Karol Gregor, Rishabh Kabra, Sébastien Racanière, Théophane Weber, David Raposo, Adam Santoro, Laurent Orseau, Tom Eccles, et al. An investigation of model-free planning. In *International Conference on Machine Learning*, pages 2464–2473. PMLR, 2019.
- [14] Demis Hassabis, Dhharshan Kumaran, Christopher Summerfield, and Matthew Botvinick. Neuroscience-inspired artificial intelligence. *Neuron*, 95(2):245–258, 2017.
- [15] Jeffrey R Hollerman, Leon Tremblay, and Wolfram Schultz. Involvement of basal ganglia and orbitofrontal cortex in goal-directed behavior. *Progress in brain research*, 126:193–215, 2000.
- [16] Dinesh Jayaraman, Frederik Ebert, Alexei A. Efros, and Sergey Levine. Time-agnostic prediction: Predicting predictable video frames. In *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. OpenReview.net, 2019.

- [17] Leslie Pack Kaelbling. Learning to achieve goals. In Ruzena Bajcsy, editor, *Proceedings of the 13th International Joint Conference on Artificial Intelligence. Chambéry, France, August 28 - September 3, 1993*, pages 1094–1099. Morgan Kaufmann, 1993.
- [18] Taesup Kim, Sungjin Ahn, and Yoshua Bengio. Variational temporal abstraction. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 11566–11575, 2019.
- [19] Thanard Kurutach, Aviv Tamar, Ge Yang, Stuart J. Russell, and Pieter Abbeel. Learning plannable representations with causal infogan. In Samy Bengio, Hanna M. Wallach, Hugo Larochelle, Kristen Grauman, Nicolò Cesa-Bianchi, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 31: Annual Conference on Neural Information Processing Systems 2018, NeurIPS 2018, December 3-8, 2018, Montréal, Canada*, pages 8747–8758, 2018.
- [20] Kara Liu, Thanard Kurutach, Christine Tung, Pieter Abbeel, and Aviv Tamar. Hallucinative topological memory for zero-shot visual planning. In *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event*, volume 119 of *Proceedings of Machine Learning Research*, pages 6259–6270. PMLR, 2020.
- [21] Yinhan Liu, Jiatao Gu, Naman Goyal, Xian Li, Sergey Edunov, Marjan Ghazvininejad, Mike Lewis, and Luke Zettlemoyer. Multilingual denoising pre-training for neural machine translation. *CoRR*, abs/2001.08210, 2020.
- [22] Piotr Miłoś, Łukasz Kuciński, Konrad Czechowski, Piotr Kozakowski, and Maciek Klimek. Uncertainty-sensitive learning and planning with ensembles. *arXiv preprint arXiv:1912.09996*, 2019.
- [23] Suraj Nair and Chelsea Finn. Hierarchical foresight: Self-supervised learning of long-horizon tasks via visual subgoal generation. In *8th International Conference on Learning Representations, ICLR 2020, Addis Ababa, Ethiopia, April 26-30, 2020*. OpenReview.net, 2020.
- [24] Soroush Nasiriany, Vitchyr Pong, Steven Lin, and Sergey Levine. Planning with goal-conditioned policies. In Hanna M. Wallach, Hugo Larochelle, Alina Beygelzimer, Florence d’Alché-Buc, Emily B. Fox, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 32: Annual Conference on Neural Information Processing Systems 2019, NeurIPS 2019, December 8-14, 2019, Vancouver, BC, Canada*, pages 14814–14825, 2019.
- [25] Giambattista Parascandolo, Lars Buesing, Josh Merel, Leonard Hasenclever, John Aslanides, Jessica B Hamrick, Nicolas Heess, Alexander Neitz, and Theophane Weber. Divide-and-conquer monte carlo tree search for goal-directed planning. *arXiv preprint arXiv:2004.11410*, 2020.
- [26] Karl Pertsch, Oleh Rybkin, Frederik Ebert, Shenghao Zhou, Dinesh Jayaraman, Chelsea Finn, and Sergey Levine. Long-horizon visual planning with goal-conditioned hierarchical predictors. In Hugo Larochelle, Marc’Aurelio Ranzato, Raia Hadsell, Maria-Florina Balcan, and Hsuan-Tien Lin, editors, *Advances in Neural Information Processing Systems 33: Annual Conference on Neural Information Processing Systems 2020, NeurIPS 2020, December 6-12, 2020, virtual*, 2020.
- [27] Karl Pertsch, Oleh Rybkin, Jingyun Yang, Shenghao Zhou, Konstantinos G. Derpanis, Kostas Daniilidis, Joseph J. Lim, and Andrew Jaegle. Keyframing the future: Keyframe discovery for visual prediction and planning. In Alexandre M. Bayen, Ali Jadbabaie, George J. Pappas, Pablo A. Parrilo, Benjamin Recht, Claire J. Tomlin, and Melanie N. Zeilinger, editors, *Proceedings of the 2nd Annual Conference on Learning for Dynamics and Control, L4DC 2020, Online Event, Berkeley, CA, USA, 11-12 June 2020*, volume 120 of *Proceedings of Machine Learning Research*, pages 969–979. PMLR, 2020.
- [28] Silviu Pitis, Harris Chan, Stephen Zhao, Bradley C. Stadie, and Jimmy Ba. Maximum entropy gain exploration for long horizon multi-goal reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning, ICML 2020, 13-18 July 2020, Virtual Event*, volume 119 of *Proceedings of Machine Learning Research*, pages 7750–7761. PMLR, 2020.

- [29] Stanislas Polu and Ilya Sutskever. Generative language modeling for automated theorem proving. *arXiv preprint arXiv:2009.03393*, 2020.
- [30] Stuart Russell and Peter Norvig. Artificial intelligence: A modern approach. ed. 3. 2010.
- [31] Nikolay Savinov, Alexey Dosovitskiy, and Vladlen Koltun. Semi-parametric topological memory for navigation. In *6th International Conference on Learning Representations, ICLR 2018, Vancouver, BC, Canada, April 30 - May 3, 2018, Conference Track Proceedings*. OpenReview.net, 2018.
- [32] Johann M Schumann. *Automated theorem proving in software engineering*. Springer Science & Business Media, 2001.
- [33] David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, et al. Mastering the game of go with deep neural networks and tree search. *nature*, 529(7587):484–489, 2016.
- [34] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Thore Graepel, Timothy Lillicrap, Karen Simonyan, and Demis Hassabis. A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 1144:1140–1144, 2018.
- [35] David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, Timothy P. Lillicrap, Karen Simonyan, and Demis Hassabis. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. *ArXiv*, abs/1712.01815, 2017.
- [36] David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George Van Den Driessche, Thore Graepel, and Demis Hassabis. Mastering the game of Go without human knowledge. *Nature*, 2017.
- [37] Gregory J. Stein, Christopher Bradley, and Nicholas Roy. Learning over subgoals for efficient navigation of structured, unknown environments. In *2nd Annual Conference on Robot Learning, CoRL 2018, Zürich, Switzerland, 29-31 October 2018, Proceedings*, volume 87 of *Proceedings of Machine Learning Research*, pages 213–222. PMLR, 2018.
- [38] Richard S. Sutton, Joseph Modayil, Michael Delp, Thomas Degris, Patrick M. Pilarski, Adam White, and Doina Precup. Horde: a scalable real-time architecture for learning knowledge from unsupervised sensorimotor interaction. In Liz Sonenberg, Peter Stone, Kagan Tumer, and Pinar Yolum, editors, *10th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2011), Taipei, Taiwan, May 2-6, 2011, Volume 1-3*, pages 761–768. IFAAMAS, 2011.
- [39] Richard S. Sutton, Doina Precup, and Satinder P. Singh. Between mdps and semi-mdps: A framework for temporal abstraction in reinforcement learning. *Artif. Intell.*, 112(1-2):181–211, 1999.
- [40] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 5998–6008, 2017.
- [41] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In Isabelle Guyon, Ulrike von Luxburg, Samy Bengio, Hanna M. Wallach, Rob Fergus, S. V. N. Vishwanathan, and Roman Garnett, editors, *Advances in Neural Information Processing Systems 30: Annual Conference on Neural Information Processing Systems 2017, December 4-9, 2017, Long Beach, CA, USA*, pages 5998–6008, 2017.
- [42] Alexander Sasha Vezhnevets, Simon Osindero, Tom Schaul, Nicolas Manfred Otto Heess, Max Jaderberg, David Silver, and Koray Kavukcuoglu. Feudal networks for hierarchical reinforcement learning. *ArXiv*, abs/1703.01161, 2017.

- [43] Oriol Vinyals, Igor Babuschkin, Wojciech M. Czarnecki, Michaël Mathieu, Andrew Dudzik, Junyoung Chung, David H. Choi, Richard Powell, Timo Ewalds, Petko Georgiev, Junhyuk Oh, Dan Horgan, Manuel Kroiss, Ivo Danihelka, Aja Huang, Laurent Sifre, Trevor Cai, John P. Agapiou, Max Jaderberg, Alexander S. Vezhnevets, Rémi Leblond, Tobias Pohlen, Valentin Dalibard, David Budden, Yury Sulsky, James Molloy, Tom L. Paine, Caglar Gulcehre, Ziyu Wang, Tobias Pfaff, Yuhuai Wu, Roman Ring, Dani Yogatama, Dario Wünsch, Katrina McKinney, Oliver Smith, Tom Schaul, Timothy Lillicrap, Koray Kavukcuoglu, Demis Hassabis, Chris Apps, and David Silver. Grandmaster level in starcraft ii using multi-agent reinforcement learning. *Nature*, 575(7782):350–354, Nov 2019.
- [44] David Wilkins. Using patterns and plans in chess. *Artif. Intell.*, 14(2):165–203, 1980.
- [45] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, and Jamie Brew. Huggingface’s transformers: State-of-the-art natural language processing. *CoRR*, abs/1910.03771, 2019.
- [46] Yuhuai Wu, Albert Jiang, Jimmy Ba, and Roger Grosse. Int: An inequality benchmark for evaluating generalization in theorem proving. *arXiv preprint arXiv:2007.02924*, 2020.
- [47] Lunjun Zhang, Ge Yang, and Bradly C. Stadie. World model as a graph: Learning latent landmarks for planning. *CoRR*, abs/2011.12491, 2020.

Checklist

1. For all authors...
 - (a) Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope? [\[Yes\]](#)
 - (b) Did you describe the limitations of your work? [\[Yes\]](#) In the Section 5
 - (c) Did you discuss any potential negative societal impacts of your work? [\[Yes\]](#) In Section 6.1
 - (d) Have you read the ethics review guidelines and ensured that your paper conforms to them? [\[Yes\]](#)
2. If you are including theoretical results...
 - (a) Did you state the full set of assumptions of all theoretical results? [\[N/A\]](#)
 - (b) Did you include complete proofs of all theoretical results? [\[N/A\]](#)
3. If you ran experiments...
 - (a) Did you include the code, data, and instructions needed to reproduce the main experimental results (either in the supplemental material or as a URL)? [\[Yes\]](#) We include the URL to the code repository in Section 1.
 - (b) Did you specify all the training details (e.g., data splits, hyperparameters, how they were chosen)? [\[Yes\]](#) See Appendix sections B, D and C.
 - (c) Did you report error bars (e.g., with respect to the random seed after running experiments multiple times)? [\[Yes\]](#)
 - (d) Did you include the total amount of compute and the type of resources used (e.g., type of GPUs, internal cluster, or cloud provider)? [\[Yes\]](#) In Appendix K
4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets...
 - (a) If your work uses existing assets, did you cite the creators? [\[Yes\]](#) For one of domains we generate expert data with implementation of MCTS agent published with work [22], as noted in Section 4.2.
 - (b) Did you mention the license of the assets? [\[N/A\]](#) The implementation noted above is made available on github, without any licence.
 - (c) Did you include any new assets either in the supplemental material or as a URL? [\[Yes\]](#) We release our code, the URL is in Section 1.
 - (d) Did you discuss whether and how consent was obtained from people whose data you’re using/curating? [\[N/A\]](#)

- (e) Did you discuss whether the data you are using/curating contains personally identifiable information or offensive content? [N/A]
- 5. If you used crowdsourcing or conducted research with human subjects...
 - (a) Did you include the full text of instructions given to participants and screenshots, if applicable? [N/A]
 - (b) Did you describe any potential participant risks, with links to Institutional Review Board (IRB) approvals, if applicable? [N/A]
 - (c) Did you include the estimated hourly wage paid to participants and the total amount spent on participant compensation? [N/A]

A MCTS

A.1 MCTS-kSubS algorithm

In Algorithm 4 we present a general MCTS solver based on AlphaZero. Solver repeatedly queries the planner for a list of actions and executes them one by one. Baseline planner returns only a single action at a time, whereas MCTS-kSubS gives around k actions – to reach the desired subgoal (number of actions depends on a subgoal distance, which not always equals k in practice).

MCTS-kSubS operates on a high-level subgoal graph: nodes are subgoals proposed by the generator (see Algorithm 3) and edges – lists of actions informing how to move from one subgoal to another (computed by the low-level conditional policy in Algorithm 2). The graph structure is represented by *tree* variable. For every subgoal, it keeps up to C_3 best nearby subgoals (according to generator scores) along with a mentioned list of actions and sum of rewards to obtain while moving from the parent to the child subgoal.

Most of MCTS implementation is shared between MCTS-kSubS and AlphaZero baseline, as we can treat the behavioral-cloning policy as a subgoal generator with $k = 1$. All the differences between MCTS-kSubS and the baseline are encapsulated in GEN_CHILDREN function (Algorithms 5 and 6). To generate children subgoals MCTS-kSubS runs subgoal generator and low-level conditional policy, whereas the baseline uses behavioral cloning policy for that purpose.

Algorithm 4 MCTS solver (common for AlphaZero baseline and MCTS-kSubS)

Require:	L_a	action limit	function SELECT(state)
	L_p	planner calls limit	$s \leftarrow \text{state}$
	P	planning passes	path $\leftarrow []$
	γ	discount factor	while s belongs to <i>tree</i> do
	c_{puct}	exploration weight	$i \leftarrow \text{SELECT_CHILD}(s)$
	τ	sampling temperature	$s', r, \text{actions} \leftarrow \text{tree}[s][i]$
	V	value function	path.APPEND((s, i, r))
	<i>env</i>	environment	$s \leftarrow s'$
	M	environment model	return path, s
Use:	<i>tree</i>	tree structure	function EXPAND(leaf)
	$N(s, i)$	visit count	children, probs $\leftarrow \text{GEN_CHILDREN}(\text{leaf})$
	$W(s, i)$	total child-value	<i>tree</i> [leaf] \leftarrow children
	$Q(s, i)$	mean child-value	$\pi(\text{leaf}, \cdot) \leftarrow \text{probs}$
	π_e	exploration policy	for $i \leftarrow 1$ to children.LENGTH() do
# Initialize N, W, Q to zero			$s', r, \text{actions} \leftarrow \text{tree}[\text{leaf}][i]$
function SOLVER			$W(\text{leaf}, i) \leftarrow r + \gamma * V(s')$
	$s \leftarrow \text{env.RESET}()$		$N(\text{leaf}, i) \leftarrow 1$
	solution $\leftarrow []$	▷ List of actions	$Q(\text{leaf}, i) \leftarrow W(\text{leaf}, i)$
	for $1 \dots L_p$ do		function UPDATE(path, leaf)
	actions $\leftarrow \text{PLANNER}(s)$		quality $\leftarrow V(\text{leaf})$
	for a in actions do		for $s, i, r \leftarrow \text{reversed}(\text{path})$ do
	$s', r \leftarrow \text{env.STEP}(a)$		quality $\leftarrow r + \gamma * \text{quality}$
	solution.APPEND(a)		$W(s, i) \leftarrow W(s, i) + \text{quality}$
	$s \leftarrow s'$		$N(s, i) \leftarrow N(s, i) + 1$
	if solution.LENGTH() $> L_a$ then		$Q(s, i) \leftarrow \frac{W(s, i)}{N(s, i)}$
	return None		function SELECT_CHILD(s)
	if <i>env.SOLVED</i> () then		$U(s, i) \leftarrow \sqrt{\sum_{i'} N(s, i') / (1 + N(s, i))}$
	return solution		$i \leftarrow \text{argmax}_i (Q(s, i) + c_{puct} \pi_e(s, i) U(s, i))$
	return None		return i
	function PLANNER(state)		function CHOOSE_ACTIONS(s)
	for $1 \dots P$ do		$i \sim \text{softmax}(\frac{1}{\tau} \log N(s, \cdot))$
	path, leaf $\leftarrow \text{SELECT}(\text{state})$		$s', r, \text{actions} \leftarrow \text{tree}[s][i]$
	EXPAND(leaf)		return actions
	UPDATE(path, leaf)		
	return CHOOSE_ACTIONS(state)		

Algorithm 5 GEN_CHILDREN for MCTS-kSubS **Algorithm 6** GEN_CHILDREN for AlphaZero

For functions *GET_PATH* and *SUB_GENERATE* see Algorithms 2 and 3.

```

function GEN_CHILDREN(state)
   $s \leftarrow \text{state}$ 
  children  $\leftarrow []$ 
  probs  $\leftarrow []$ 
  for subgoal, prob  $\leftarrow \text{SUB\_GENERATE}(s)$  do
    actions  $\leftarrow \text{GET\_PATH}(s, \text{subgoal})$ 
    if actions.EMPTY() then continue
     $r \leftarrow M.\text{REWARD\_SUM}(s, \text{actions})$ 
    children.APPEND((subgoal,  $r$ , actions))
    probs.APPEND(prob)
  return children, probs

```

Require: π_b behavioral cloning policy

```

function GEN_CHILDREN(state)
   $s \leftarrow \text{state}$ 
  children  $\leftarrow []$ 
  probs  $\leftarrow []$ 
  for  $a, \text{prob} \leftarrow \pi_b.\text{GEN\_ACTIONS}(s)$  do
     $s', r \leftarrow M.\text{NEXT\_STATE\_REWARD}(s, a)$ 
    children.APPEND(( $s', r, [a]$ ))
    probs.APPEND(prob)
  return children, probs

```

Variables $tree, N, W, Q, \pi_e$ are reused across subsequent planner invocations within a single solver run. We limit the number of planner calls L_p for better control over the computational budget for MCTS-kSubS. For MCTS pseudocode we assume a slightly modified version of SUB_GENERATE function (defined originally in Algorithm 3). We presume that the function along with subgoals returns also their respective probabilities – as MCTS needs them to guide exploration.

A.2 Detailed results of MCTS-kSubS

We evaluate MCTS-based approaches on INT proofs of length 15. We set $c_{puct} = \tau = 1$ and $\gamma = 0.99$. We tuned P on MCTS (AlphaZero) baseline and we run three variants with $P \in \{5, 15, 50\}$. We run MCTS-kSubS ($k = 3$) with the same set of parameters and with kSubS-specific parameters fixed to $C_2 = C_3 = 4$ (in order to match the setup for corresponding INT BF-kSubS experiments).

We limit the maximum number of actions to $L_a = 24$ for both methods. Having the same number of planning passes P , during a single call MCTS-kSubS visits k -times more new states than the baseline (because of states visited by the low-level conditional policy). Therefore, to ensure a similar computational budget, we limit the number of planner calls to $L_p = 8$ for MCTS-kSubS and to $L_p = 24$ for the baseline – so the number of states visited over the course of a single solver run is similar for both methods.

Top-left part of Figure 1 illustrates results of MCTS experiments. For every number of planning passes P , MCTS-kSubS has significantly higher success rate than the corresponding baseline experiment. The highest difference is 0.52 for $P = 5$ (0.88 for MCTS-kSubS, 0.36 for the baseline) and slightly decreases with increasing number of passes to still impressive 0.43 for $P = 50$ (0.91 for MCTS-kSubS, 0.48 for the baseline). Comparing MCTS-kSubS for $P = 5$ with the baseline for $P = 50$, shows advantage of our method still by a significant margin of 0.40, despite having 10 times smaller computational budget.

MCTS-kSubS performed better also in terms of run time. For every tested P it was at least twice as fast as the corresponding baseline experiment.

High effectiveness of MCTS-kSubS, in terms of both search success rate as well as run time, shows that our kSubS method is not specific to BestFS planner, but potentially allows to boost a wide range of other planners.

B Architectures and hyperparameters

B.1 Transformer

For INT and Rubik we use mBART [21] – one of the state-of-the-art sequence-to-sequence transformer architectures. To speed up training and inference we use its lightweight version. We reduced the dimensionality of the model, so the number of learned parameters decreased from the original 680M to 45M. The set of our hyperparameters matches the values proposed in [40]: we used 6 layers of encoder and 6 layers of decoder; we adjusted model’s dimension to 512 and number of attention

heads to 8; the inner-layer of position-wise fully connected networks had dimensionality 2048. The difference in our model’s size compared to 65M parameters reported in [40] results from vocabulary size. For our problems, it is enough to have 10-70 distinct tokens, whereas natural language models require a much larger vocabulary (tens of thousands of tokens).

B.2 Sokoban

In Sokoban, we use only two neural network architectures: one for generating subgoals, and one for assigning value.

We took the value function network from the training of an expert agent [22]. For the subgoal generator network we used the same convolutional architecture, with two expectations. First, instead of predicting single regression target we predicted distribution over $d \times d \times 7 + 1$ classes. Secondly, we added batch norm layers between convolutional layers to speed up training. To make the comparison between BestFS and BF-kSubS fair, we also evaluated the training of expert from [22] with additional batch norm layers, but it turned out to actually hurt the performance of the expert.

C Data processing and datasets

C.1 Sokoban

Dataset. We collected expert datasets using an RL agent (MCTS-based) from [22]. Precisely, we trained 3 agents on Sokoban boards of different sizes (12×12 , 16×16 and 20×20 , all with four boxes). During the training process, we collected all successful trajectories, in the form of sequences of consecutive states. The number of trajectories in our datasets were: 154000 for 12×12 boards, 45000 for 16×16 boards and 21500 for 20×20 boards. The difference comes from the fact that the larger boards take more time to solve, hence fewer data is collected in the same time span.

Value. As a byproduct of the training mentioned in the previous paragraph, we got the trained value network used by the agent. The target for learning values was defined as a discounted distance to the solution (with the discount factor of 0.99). We utilized this network as a value function in our algorithms.

Subgoal generation. For a given state the generation of subgoals is depicted in Algorithm 7. We maintain a queue of modified states (MS). Iteratively we take a MS from queue, concatenate it with state and pass through subgoal generator network (`subgoal_net.SORTED_PREDICTIONS`). This produces a probability distribution over candidates for further modifications of given MS. We take the most probable candidates, apply each of them to MS, and add the new modified states to the queue. If among the best subgoal generator network predictions there is a special "valid subgoal" token (encoded with $d \times d \times 7 + 1$), we put MS to subgoal candidates list (`subgoals_and_probs`). During this process, each MS is assigned the probability, which is a product of probabilities of modifications, leading to this MS. When the queue is empty, we take subgoal candidates and choose the ones with the highest probability such that the target probability (C_5) is reached (similar to Algorithm 3). The generation of subgoals for a given state is illustrated in Figure 5.

This process is designed to be computationally efficient. The majority of subgoals differ from the input by only several pixels, which leads to short paths of point-wise modifications. Note, that we do not utilize any Sokoban-specific assumptions.

Datapoints for training. Preparing data points for the training of the generator is described in Algorithm 8. For each trajectory in the dataset, we choose randomly 10% of state pairs for the training (we do not use all states from a trajectory in order to reduce the correlation in the data).

Performance of RL agent. We observed that each of the three RL agents we used (for 12×12 , 16×16 and 20×20 boards), had a significantly lower success rate than our method’s counterparts (that learns from these agents). For 12×12 boards it could solve around 78% of problems, for 16×16 boards it dropped to 67% and for 20×20 it was only 60%.

Algorithm 7 Sokoban subgoal generator

Require: d dimension of a board
 $internal_cl$ a number between 0 and 1
 $subgoal_net$ CNN returning distribution over modifications.

function GENERATE_SUBGOALS($state$)
 $subgoals_and_probs \leftarrow []$
 $q \leftarrow Queue()$ ▷ FIFO queue
 $q.INSET((state, 1))$
 while not q not empty **do**
 $modified_state, parent_prob \leftarrow q.POP()$
 $network_input \leftarrow CONCATENATE(state, modified_state)$
 $predictions, probs \leftarrow subgoal_net.SORTED_PREDICTIONS(network_input)$
 $total_p \leftarrow 0$
 for $prediction, p \in (predictions, probs)$ **do**
 if $total_p \geq internal_cl$ **then break**
 $total_p \leftarrow total_p + p$
 if $prediction = d \times d \times 7 + 1$ **then**
 $subgoals_and_probs.ADD((modified_state, parent_prob \times p))$
 else
 $new_modified_state \leftarrow APPLY_CHANGE(modified_state, prediction)$
 $q.INSET((new_modified_state, parent_prob \times p))$
 $subgoals \leftarrow SORT_BY_PROBABILITY(subgoals)$
 $total_p \leftarrow 0$
 for $subgoal, p \in subgoals_and_probs$ **do**
 if $total_p > C_5$ **then break**
 $subgoals.ADD(state)$
 $total_p \leftarrow total_p + p$
 return $subgoals$

function APPLY_CHANGE($state, modification$)
 ▷ $modification$ is an integer in range $[1, d \times d \times 7]$ encoding which pixel of $state$
 ▷ to change (and to which value).
 $row \leftarrow \frac{modification}{d \times 7}$ ▷ Integer division
 $column \leftarrow \frac{modification - row \times d \times 7}{7}$ ▷ Integer division
 $depth \leftarrow modification - row \times d \times 7 - column \times 7$
 $modified_state \leftarrow state$
 $SET_TO_ZEROES(modified_state[row, column])$
 $modified_state[row, column, depth] \leftarrow 1$
 return $modified_state$

Algorithm 8 Sokoban generate network inputs and targets

Require: d dimension of a board

```
function GENERATE_INPUTS_AND_TARGETS(state, subgoal)
    inputs  $\leftarrow$  [] ▷ empty list
    targets  $\leftarrow$  [] ▷ empty list
    modified_state  $\leftarrow$  state
    input  $\leftarrow$  CONCATENATE(state, modified_state)
    inputs.APPEND(input)
    target_class_num  $\leftarrow$  0
    for  $i \in 1 \dots d$  do
        for  $j \in 1 \dots d$  do
            for  $c \in 1 \dots 7$  do
                target_class_num  $\leftarrow$  target_class_num + 1
                if subgoal[i, j, c] = 1 AND modified_state[i, j, c] = 0 then
                    targets.APPEND(target_class_num)
                    ▷ Numpy notation, replace pixel values on position (i, j) with values
                    ▷ from subgoal
                    modified_state[i, j, :]  $\leftarrow$  subgoal[i, j, :]
                    input  $\leftarrow$  CONCATENATE(state, modified_state)
                    inputs.APPEND(input)
            ▷ Last target: no more changes to the modified_state are needed (class enumerated
            ▷ with  $d \times d \times 7 + 1$ )
            targets.APPEND( $d \times d \times 7 + 1$ )
    return inputs, targets
```

C.2 INT

State representation. A state in INT consists of objectives to prove and ground truth assumptions, which are logic statements that are assumed to hold and can be used in proving. Each objective, as well as each ground truth assumption, is a mathematical statement. In our setup, as in the original paper [46], there is always only one objective to prove, but there may be a varying number of ground truth statements.

Each mathematical statement can be converted to a string by a method `logic_statement_to_seq_string` in INT library. In our code, we represent the full state as a string in the following form (all symbols are tokens, not logical operations):

#[objective]&[1st ground truth]&[2nd ground truth]&... &[k - th ground truth]\$

For example, the state representation could look like this:

#(((b+b)*((b+b)*(b+b)))*(((b+b)+f)*(b+b))*(b+b))) ≥ 0 & (b+f) = b & (b+f) ≥ 0 \$

Action representation. An action in INT consists of a chosen axiom (one from the set of ordered field axioms, see [46, Appendix C]) and a sequence of entities onto which the axiom will be applied. An entity is an algebraic expression that is a part of the larger statement. For example, $(a + b)$ is an entity, which is a part of $(a + b) \cdot c = (1 + f)$. In our code, we represent entities by indicating the symbol of their mathematical operation, or if the entity is atomic (a single variable), by indicating the variable itself. More precisely, directly after the symbol of operation, we add a special character '~'. For example, we indicate $(a + b)$ inside $(a + b) \cdot c$ in the following way: $(a + \sim b) \cdot c = (1 + f)$. Typically, in a logic statement there may be several entities that have the same text form, but are located in different positions, for example $(a + b)$ appears twice in $(a + b) \cdot (a + b) = (1 + 0)$. Our way of encoding actions unambiguously identifies every entity. If the action has more than one input, we use more different indicators.

Low-level conditional policy input representation. Low-level conditional policy takes as an input two states: s and s' , where s is the initial state and s' is the subgoal. The input is constructed in the following way: first, we represent both s and s' as strings and then we find the difference (character

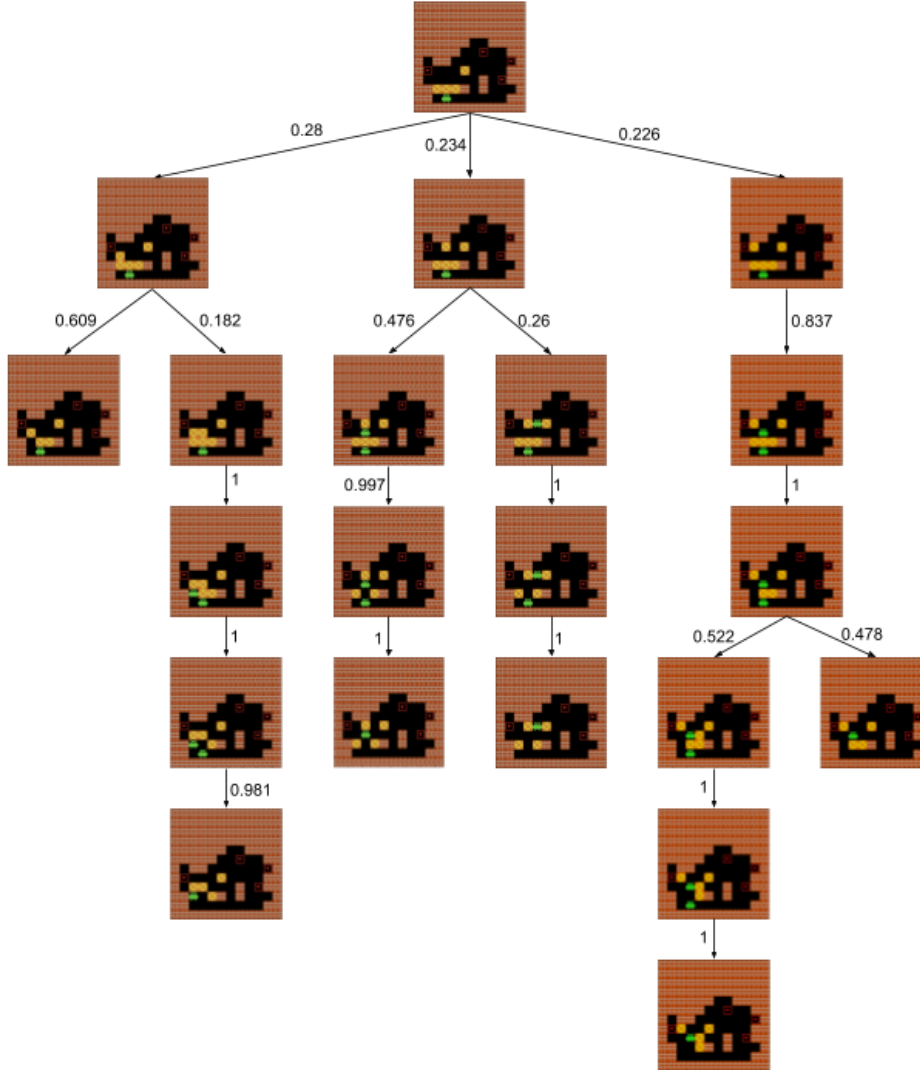


Figure 5: A detailed view of subgoal generation for Sokoban. Arrow represent probabilities of a given modification. Final subgoals are located in the leaves.

delta) between these strings using the function `ndiff` from `diffli`⁵ Python library. We observed that using the character delta, instead of the concatenation of s and s' , significantly improved the performance.

C.3 Rubik's Cube

State representation The state of the Rubik's Cube is determined by the arrangement of 54 colored labels on its faces. Therefore, to represent the observations we simply put the labels in a fixed order. An example state is as follows:

?byywygrygobbrboorgwbowryooogywbgywrrroyogyowwbrwbbgg\$,

where the tokens b, g, o, r, w, y stand for *blue, green, orange, red, white, and yellow*. The consecutive blocks of 9 tokens correspond to consecutive faces of the cube. Observe, that not every permutation of colors is valid. For example, the tokens on positions 5, 14, 23, 32, 41, and 50 correspond to centers of faces, thus they are fixed. There are more such constraints, but they are irrelevant to the pipeline itself.

⁵<https://docs.python.org/3/library/difflib.html>

Action representation In our experiments we use quarter turns, i.e. an action corresponds to rotating a face by 90° , either clockwise or counterclockwise. Since the action space contains only 12 elements, we use unique tokens to represent each of them.

Low-level conditional policy input representation. The conditional policy takes two states s and s' , which correspond to the current state and the state to be reached. To represent such pairs, on every position we put a token corresponding to a pair of colors – one located on that position in s and the other in s' . Since there are only 6 distinct colors on the Rubik’s Cube, this requires using only 36 tokens.

D Training details

D.1 INT and Rubik’s Cube

D.1.1 Transformer training

For transformer training and inference we used HuggingFace’s Transformers library [45]. We did not use any pretrained checkpoints from HuggingFace model hub. We took mBART model class instead – and trained it from scratch in a supervised way using HuggingFace’s training pipeline. We generated (or loaded from disk) a fresh dataset for every epoch. Training batch was of size 32. For regularization, we set $dropout = 0.1$, but we did not use label smoothing (as opposed to [40]).

For the Adam optimizer we set $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. Learning rate schedule followed the formula:

$$lr = peak_lr * \min\left(\frac{step_num}{warmup_steps}, \sqrt{\frac{warmup_steps}{step_num}}\right),$$

where $peak_lr = 3 \cdot 10^{-4}$ and $warmup_steps = 4000$.

The schedule curve matches the one proposed in [40], but they use $peak_lr \approx 7 \cdot 10^{-4}$.

D.1.2 Sequence generation

We used beam search with the number of beams set to 16 for INT and to 32 for Rubik’s Cube. The number of returned sequences varied from 1 to 4 depending on the network type.

We set softmax temperature to 1 by default. For Rubik’s Cube subgoal generator we tuned this parameter and the value of 0.5 performed best. We conducted a corresponding experiment for the baseline policy, but the temperature did not impact results in this case, as the policy outputs only a single token. For INT we did not tune the temperature, so we kept the value of 1.

D.1.3 Random seeds

Due to use of the supervised learning, we observed little variance with respect to the random initialization. We tested this for the subgoal generator on proofs of length 10 and for $k = 3$. Namely, we trained 5 models of the subgoal generator, starting from different initializations. The success rate barely varied, as they stayed in the interval $[0.990, 0.992]$. In the other experiments, we used a single seed.

D.2 Sokoban

For training of the convolutional networks in Sokoban we set the learning rate to 10^{-4} and the number of epochs to 200.

E Wall-time for kSubS

As indicated in Table 2, kSubS builds smaller search graphs. This has the practical advantage of making fewer neural network calls and consequently a substantially better wall-time.

The gains might be as high as 7 times due to costly sequential calls of transformer networks, see Table 4.

Proof length	5		10		15	
Method	BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)	BestFS	BF-kSubS (ours)
Total wall-time	4h 12m	3h 44m	29h 22m	5h 55m	69h 15m	9h 22m
Avg. generator calls	NA	3.04	NA	3.89	NA	6.23
Avg. value calls	23.34	4.01	112.35	4.80	159.46	6.59
Avg. policy calls	22.41	8.43	112.21	13.09	161.02	20.29

Table 4: Resources consumption for INT. We present evaluation on 1000 proofs and split into calls of subgoal generator network (used only in the subgoal search), value network and policy network (we report an average number of calls for a single proof).

F Value errors

F.1 Sokoban Analysis

Here, we present details related to the Sokoban analysis from Section 4.6. We sampled 500 Sokoban boards (with dimension (12, 12)). For each board, we calculated a full state-action graph and minimal distance from each state to the solution (this was needed to compute $S(s)$ sets later on). Since Sokoban graphs can be very large we set the limit on the graph size to 200000, which left us with 119 boards. Next, for each board, we took the sequence of states from the shortest solving path from the initial state (let us call this dataset as *shortest-paths* - SP). For each pair of consecutive states in SP, we calculated the difference of the value estimation, and averaged them, which gave us a mean one-step improvement of 0.338 (95% confidence interval - [0.297, 0.380]).

l	Value decrease prob.
1	0.316
2	0.217
3	0.080
4	0.020

Table 5

To calculate the variance of value function estimates for $S(s)$ we took SP, and limit it to states s such that $|S(s)| \geq 5$ (lets denote it as SP5). We calculated variance for each $s \in SP5$ separately. This gave us a mean variance of 0.844 (95% confidence interval - [0.749, 0.939]). The same set SP5 was used to calculate probabilities related to overoptimistic errors on Sokoban described at the end of Section 4.6.

F.2 INT analysis

Due to the size of state spaces, it is impossible to search over the entire space of INT formulas to find the shortest proofs. We instead analyze value estimations along proofs generated by INT engine. The monotonicity analysis in Section 4.6 was performed using 100 such proofs of length 10. The probabilities of value decrease for different step lengths l are presented in Table 5.

G Example subgoals

G.1 Example subgoals sets

In this section, we present some example outcomes of the subgoal generator for INT and Sokoban.

G.1.1 INT

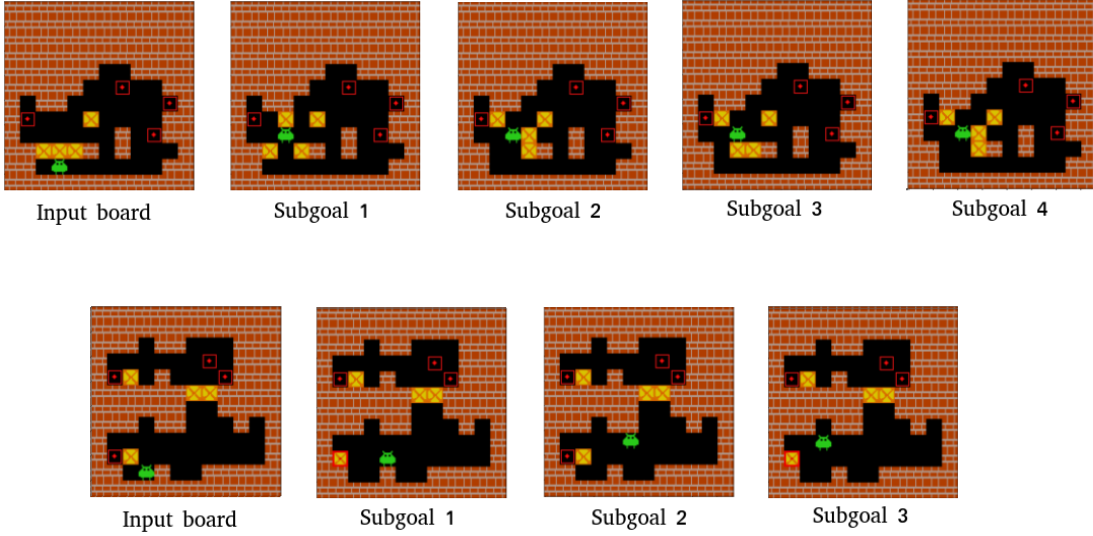
In (1) and (2) we provide two examples of applying the subgoal generator (trained on proofs of length 5) to the given states in INT. The number of subgoals varies since not all of the outputs generated by the network could be reached by the conditional low-level policy.

$$\begin{aligned}
&\text{Input state: } (((b \cdot b) + ((b + (b + f)) \cdot b)) + (f + f)) \geq (((b + (b + b)) \cdot b) + 0) + c \\
&\text{Ground truth 1: } (b + f) = b \\
&\text{Ground truth 2: } (f + f) \geq c \\
&\text{Subgoal 1: } ((b \cdot b) + ((b + (b + f)) \cdot b)) = (((b + (b + b)) \cdot b) + 0) \\
&\text{Subgoal 2: } (((b + (b + f)) \cdot b) + (b \cdot b)) = (((b + (b + b)) \cdot b) + 0) \\
&\text{Subgoal 3: } ((b \cdot b) + ((b + (f + b)) \cdot b)) = (((b + (b + b)) \cdot b) + 0)
\end{aligned} \tag{1}$$

$$\begin{aligned}
&\text{Input state: } (((((b \cdot (\frac{1}{b})) + a)^2) + (c + (\frac{1}{b}))) = ((((((b \cdot (\frac{1}{b})) \cdot b) + a) \cdot (a + 1)) + c) + (\frac{1}{b}))) \\
&\text{Subgoal 1: } (((((b \cdot (\frac{1}{b})) + a)^2) + (c + (\frac{1}{b}))) = (((((b \cdot (\frac{1}{b})) + a) \cdot (a + 1)) + c) + (\frac{1}{b}))) \\
&\text{Subgoal 2: } (((((b \cdot (\frac{1}{b})) + a) \cdot ((b \cdot (\frac{1}{b})) + a)) + (c + (\frac{1}{b}))) = (((((b \cdot (\frac{1}{b})) + a) \cdot (a + 1)) + c) + (\frac{1}{b})))
\end{aligned} \tag{2}$$

G.1.2 Sokoban

Here we present two examples of the outcomes of the subgoal generator trained for 12×12 boards:



G.2 Example solutions with subgoals

In this section, we present several examples of solutions obtained with our method. For simplicity, we only show the subgoal states on which the successful trajectories were constructed. In our setup, the last subgoal is always a solution.


G.2.1 INT

An example solution of INT problem of length 5:


$$\begin{aligned}
&\text{Problem: } (((0 \cdot ((a + 0) + (-a \cdot 1))) \cdot (\frac{1}{(0^2)})) + ((a + 0) + (0^2))) \geq (((0^2) + (1 + (a + 0))) + b) + (-((a + 0) + f)) \\
&\text{1st subgoal: } (((0 \cdot ((a + 0) + (-a \cdot 1))) \cdot (\frac{1}{(0^2)})) + ((a + 0) + (0^2))) = ((0^2) + (1 + (a + 0))) \\
&\text{2nd subgoal: } (((0 \cdot ((0 + a) + (-a \cdot 1))) \cdot (\frac{1}{(0^2)})) + ((a + 0) + (0^2))) = (1 + ((a + 0) + (0^2))) \\
&\text{3rd subgoal: } (0 + a) = (a \cdot 1) \\
&\text{4th subgoal: } a = a
\end{aligned} \tag{3}$$

G.2.2 Sokoban


An example solution of Sokoban board:




Input board




subgoal number 1




subgoal number 2




subgoal number 3




subgoal number 4




subgoal number 5




subgoal number 6




subgoal number 7



subgoal number 8



subgoal number 9



subgoal number 10

H Baselines

Our first baseline is the low-level policy trained with behavioral cloning from the expert data trained to solve the problem (contrary to the low-level conditional policy P , which aims to achieve subgoals). Such policy was used [46]. We verified that our behavioral cloning policy reproduces the results from [46] for proofs of lengths 5.

MCTS. As a baseline for MCT-kSubS we used an AlphaZero-based MCTS planner described in Appendix A.1.

BestFS. The baseline for BF-kSubS is a low-level planning. We substitute SUB_GENERATE with a function returning adjacent states indicated by the most probable actions of behavioral cloning policy, see Algorithm 9.

Algorithm 9 Low-level generator

Require: C_3 number of states to produce
 C_4 number of beams in sampling
 π_b behavioral cloning policy
 M model of the environment

function SUB_GENERATE(s)
 $\text{actions} \leftarrow \text{BEAM_SEARCH}(\pi_b, s; C_3, C_4)$
 $\text{subgoals} \leftarrow []$
 for $\text{action} \in \text{actions}$ **do**
 $s \leftarrow M.\text{NEXT_STATE}(s, \text{action})$
 $\text{subgoals}.\text{APPEND}(s)$
 return subgoals

I Simple planners

An appropriate search mechanism is an important design element of our method. To show this, we evaluate an alternative, simpler procedure used by e.g. [8] for subgoal-based planning. It works by

sampling independent sequences of subgoals and selects the best one. This method solved none of 1000 Rubik’s Cube instances despite using the same subgoal-generator as BF-kSubS (which has a success rate of 0.999 with a comparable computational budget).

J Investigation of baseline-BestFS on Rubik’s Cube

To obtain the training datasets on the Rubik’s Cube environment, we generated random paths starting from a solved cube and stored them in the reversed order. These backward solutions are highly sub-optimal: for example, the states obtained by 10 random moves are usually in a distance of about 6 - 7 steps from the final state and the gap gets much larger for a higher number of moves. This means that on collected trajectories only some of the actions indeed lead to the solution and the majority of them only introduce noise.

We claim that such noisy data is sufficient for low-level conditional policy and the generator, but not for the behavioral cloning policy, which is used by the baseline (see Appendix H).

Let us consider two states s and s' which lie (in this order) on one random trajectory τ within small distance k (e.g. $k = 4$). In most cases the fragment of τ between s and s' is the optimal path joining s with s' . Hence, for the low-level conditional policy, such trajectories are sufficient to learn how to achieve subgoals. Also, if k is large enough (say $k = 4$), it is likely that s' would be at least one step closer to the solution than s . Therefore, kSubS successfully mitigates the issue of randomness in the learning data.

However, the same does not apply to the behavioral cloning policy, which is trained to predict the expert’s action. Since it looks only one step ahead, it is much more sensitive to randomness. As noted before, in collected trajectories only a part of actions lead closer to the solution, so in many cases, the targets for training policy are highly sub-optimal and noisy, thus hard to generalize. When trained on the same data as the networks used for kSubS, the behavioral cloning policy fails to significantly exceed the performance of the random policy.

To provide a meaningful evaluation of the baseline, we also trained it on very short trajectories consisting of 10 random moves. Such a curriculum allowed the behavioral cloning policy to learn, however still suffered from randomness in data (for states located far from the solution).

We extended the training of behavioral cloning policy and did a grid-search over parameters to further improve its success rate. We managed to reach no more than 4% success rate for the best version.

Note that both kSubS and baseline policy learn from fully off-policy data. The problem of solving the Rubik’s Cube is challenging, thus learning from noisy off-policy trajectories can be simply too hard for standard algorithms.

K Technical details

K.1 Infrastructure used

We had 2 types of computational nodes at our disposal, depending on whether a job required GPU or not. GPU tasks used a single Nvidia V100 32GB card (mostly for transformer training) and Nvidia RTX 2080Ti 11GB (for evaluation) with 4 CPU cores and 16GB RAM. The typical configuration of CPU job was the Intel Xeon E5-2697 2.60GHz processor (28 cores) with 128GB memory.

Each transformer for INT was trained on a single GPU for 3 days (irrespective of proof length and the network type). INT evaluation experiments used a single GPU for a time period varying from several hours to 3 days – with baselines being the bottleneck.

Transformers for Rubik’s Cube required more training – every network was trained for 6 days on a single GPU. Because of relatively short sequences representing the cube’s state, we were able to run evaluations without GPU. We used 20 CPU cores with 20GB memory, while still fitting in a 3-day run time.

We trained and evaluated Sokoban on CPU only mostly because of rather small neural network sizes. Training time varied from 2 to 3 days, whereas evaluation took only an hour.